

ICAPS 2023 Tutorial

**Introduction to Domain  
Modeling in RDDDL  
Part 1: Language Overview**

Scott Sanner and Ayal Taitler



UNIVERSITY OF  
**TORONTO**

# Multiple Target Audiences

- ICAPS folks familiar with (P)PDDL wondering what RDDL is and when they might use it
- Planning language agnostics who are simply interested in planning for MDPs and POMDPs
- RL researchers interested in how to specify and exploit complex model structure


# RDDL Tutorial Outline

- **Part 1: Language Overview**
  - **What is probabilistic planning in PPDDL?**
  - Why do we need RDDL?
  - RDDL by example
  - Overview of RDDL solution methodologies
- **Part 2: PyRDDL Gym**

# Stochastic Domain Languages as of 2009

- **Probabilistic PDDL (PPDDL)**
  - more expressive than PSTRIPS
  - for example, *probabilistic universal* and *conditional* effects:


```
(:action put-all-blue-blocks-on-table
:parameters ( )
:precondition ( )
:effect (prob 0.9
        (forall (?b)
          (when (Blue ?b)
            (not (OnTable ?b))))))
```











- **Idea:** make some effects stochastic
- **Question:** is this sufficient to model realistic problems?

# More Realistic: Logistics?

- PPDDL Description:



*(:action load-box-on-truck-in-city*  
*:parameters* (*?b* - box *?t* - truck *?c* - city)  
*:precondition* (and (BIn *?b* *?c*) (TIn *?t* *?c*))  
*:effect* (**prob 0.7** (and (On *?b* *?t*) (not (BIn *?b* *?c*))))

- Can instantiate problems for any domain objects
  - 3 trucks:    2 planes:   3 boxes:   
- But wait... only one truck can move at a time???
- No concurrency, no time: **will FedEx care?**

# Expressivity Limitations of PPDDL

- Many PPDDL domains were tweaks of PDDL domains
  - Recipe: add success probability on some effects
    - e.g., *load-plane(p,x)* succeeds with prob 0.9
  - IPPC 2004/6, could win by determinizing / replanning
    - led to work on “probabilistically interesting” PPDDL problems (Little & Thiebaux, 2007)
- But what stochastic expressiveness is needed for modeling real-world domains?
  - Then we can ask what language is appropriate

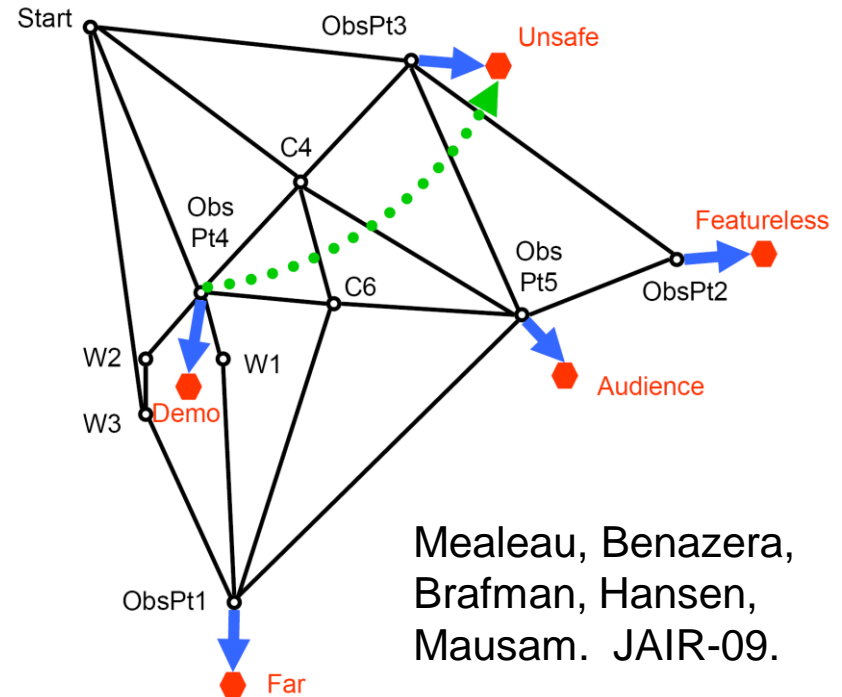
# Observation

- Planning languages direct 5+ years of research
  - PDDL and variants
  - Probabilistic PDDL (PPDDL)
- Why?
  - Domain design is time-consuming
    - So everyone (students) use existing benchmarks
  - Need for comparison
    - Planner code not always released
    - Only means of comparison is on competition benchmarks
- Implication:
  - We should choose our languages & problems well
  - Let's ask what problems we want to model / solve

What probabilistic problems  
might we want to model?




# Mars Rovers



- **Continuous**
  - Time, robot position / pose, sun angle, battery reserves...
- **Partially observable**
  - Even worse: high-dimensional partially observable

# Elevator Control

- **Concurrent Actions**
  - Elevator: up/down/stay
  - 6 elevators:  $3^6$  actions
- **Exogenous / Non-boolean**
  - Random integer arrivals (e.g., Poisson) at every floor
- **Complex Objective**
  - Minimize sum of wait times
  - Could even be nonlinear function (squared wait times)
- **Complex Action Constraints**
  - People might get annoyed if elevator reverses direction



# Traffic Control



- Concurrent
  - Multiple lights
- Continuous Variables
  - Nonlinear dynamics
- Indep. Exogenous Events
  - Multiple vehicles
- Partially observable
  - Only observe stoplines

# RDDL Tutorial Outline

- **Part 1: Language Overview**
  - What is probabilistic planning in PPDDL?
  - **Why do we need RDDL?**
  - RDDL by example
  - Overview of RDDL solution methodologies
- **Part 2: PyRDDL Gym**

# What are we missing in PPDDL?


- Independent concurrent stochastic actions & events
  - Exogenous events that scale with domain size
    - Random person arrivals at elevator floors, traffic movement
  - Resolution of stochastic or concurrent event conflicts
    - Two elevators admit passengers from same floor
  - Preconditions over joint actions (not per action)
    - Joint traffic light configurations must adhere to safety constraints
- Remedy: action-centric (P)PDDL → fluent-centric RDDDL

Need expressive decision-making  
formalism that supports complex  
stochastic **fluent** updates

Relational Dynamic Bayes Net  
+ Influence Diagram (RDDL)

a.k.a. Relational Factored MDP

# Dynamical Models & Influence Diagrams



- Dynamic Bayes Nets (DBNs) ...

- Represent state @ times  $t, t+1$

- Assume stationary distribution

- Influence Diagrams (IDs)...

- Action nodes [squares]

- Not random variables
- Rather “controlled” variables

- Utility nodes <diamonds>


- A utility conditioned on state, e.g.

$$U(X_1', X_2') = \text{if } (X_1' = X_2') \text{ then } 10 \text{ else } 0$$



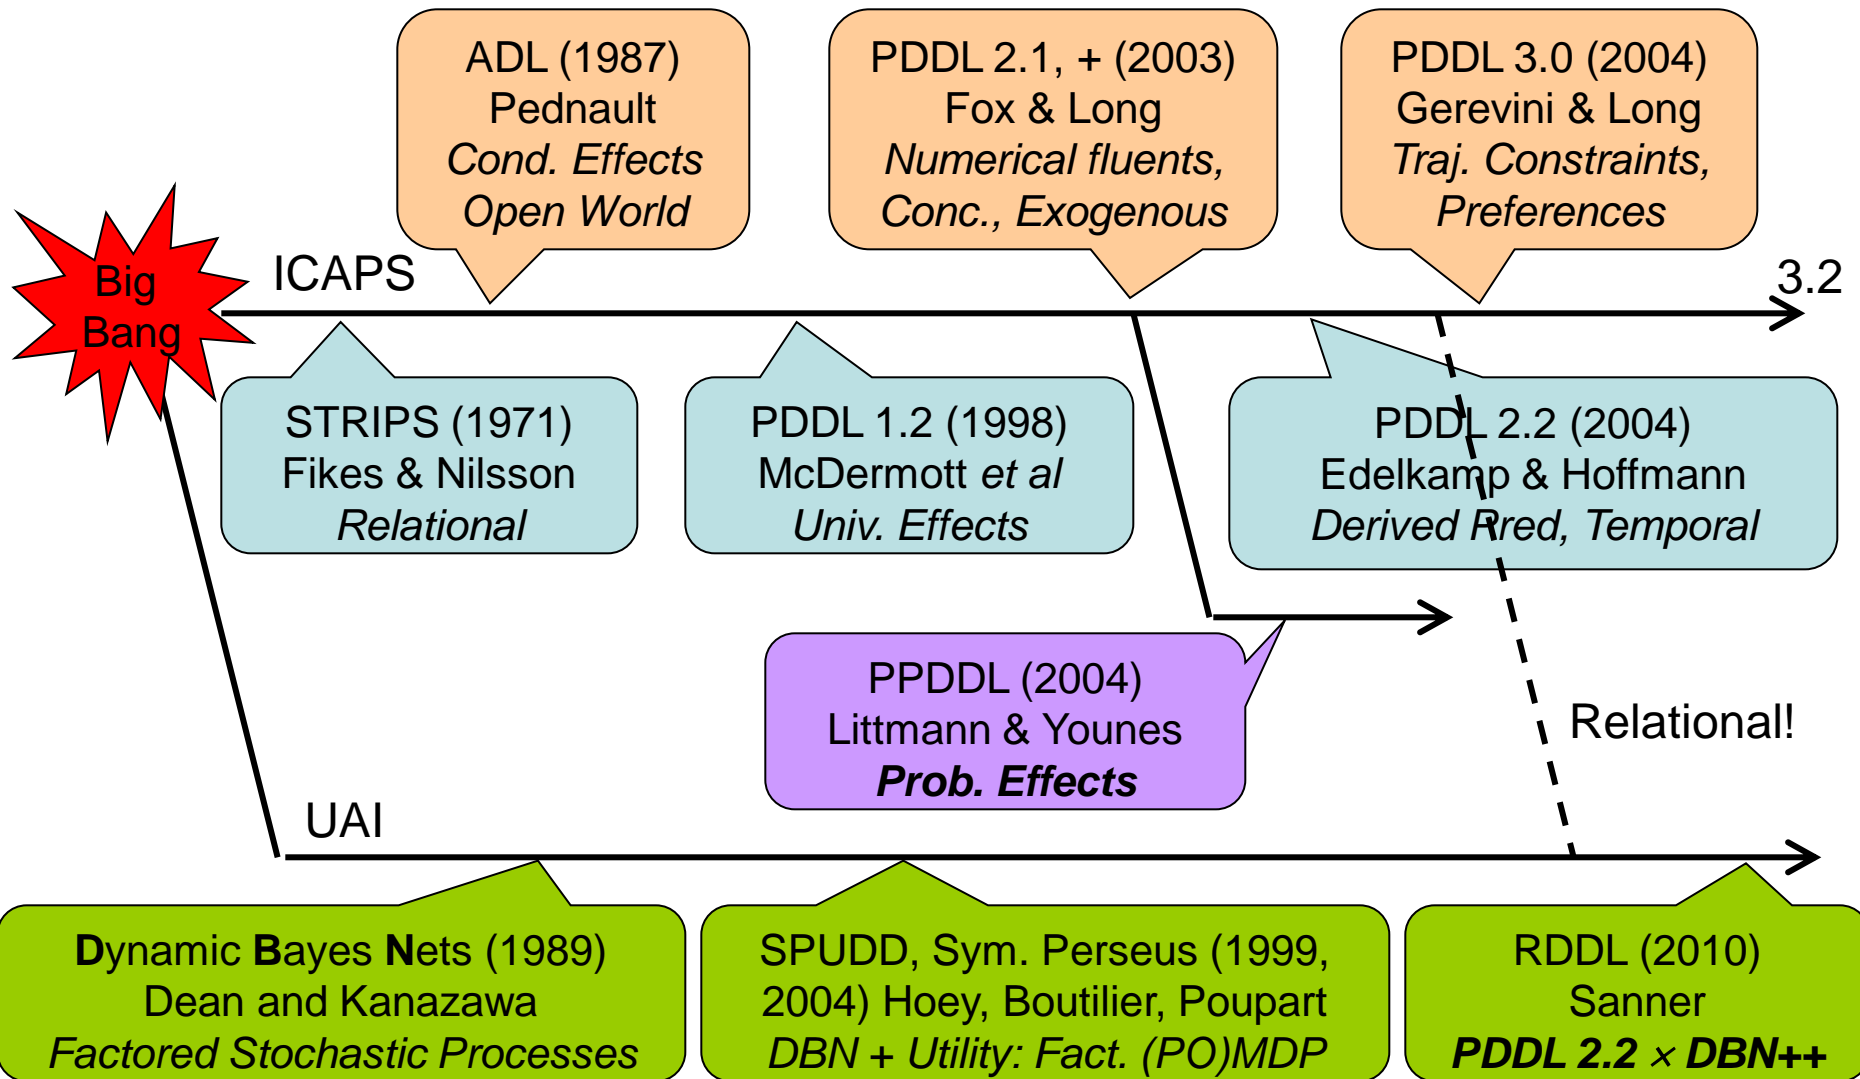
# What is RDDDL?

- Relational Dynamic Influence Diagram Language
  - Relational  
[DBN + Influence Diagram]
- Think of it as a Relational Factored (PO)MDP
  - Fluent updates are probabilistic programs





# A Brief History of (ICAPS) Time




# RDDL Tutorial Outline

- **Part 1: Language Overview**
  - What is probabilistic planning in PPDDL?
  - Why do we need RDDL?
  - **RDDL by example**
  - Overview of RDDL solution methodologies
- **Part 2: PyRDDL Gym**

Example:

How to specify a problem in  
RDDDL (that cannot be  
expressed in PPDDL)

# Wildfire Domain



- Contributed by Zhenyu Yu (School of Economics and Management, Tongji University)
  - Karafyllidis, I., & Thanailakis, A. (1997). *A model for predicting forest fire spreading using gridular automata*. Ecological Modelling, 99(1), 87-97.

# Wildfire in RDDDL

Each cell may independently stochastically ignite

```
cpfs {  
  
  burning'(?x, ?y) =  
    if ( put-out(?x, ?y) )  
      then false  
    else if (~out-of-fuel(?x, ?y) ^ ~burning(?x, ?y))  
      then Bernoulli( 1.0 / (1.0 + exp[4.5 - (sum_{?x2: x_pos, ?y2: y_pos}  
        (NEIGHBOR(?x, ?y, ?x2, ?y2) ^ burning(?x2, ?y2))]) )  
    else  
      burning(?x, ?y); // State persists  
  
  out-of-fuel'(?x, ?y) = out-of-fuel(?x, ?y) | burning(?x, ?y);  
  
};  
  
reward =  
  [sum_{?x: x_pos, ?y: y_pos} [ COST_CUTOUT*cut-out(?x, ?y) ]]  
+ [sum_{?x: x_pos, ?y: y_pos} [ COST_PUTOUT*put-out(?x, ?y) ]]  
+ [sum_{?x: x_pos, ?y: y_pos} [ COST_NONTARGET_BURN*[ burning(?x, ?y) ^ ~TARGET(?x, ?y) ]]]  
+ [sum_{?x: x_pos, ?y: y_pos}  
  [ COST_TARGET_BURN*[ (burning(?x, ?y) | out-of-fuel(?x, ?y)) ^ TARGET(?x, ?y) ]]];
```

# Power of Lifting

Simple domains  
can generate  
complex DBNs!

non-fluents game2x2 {

domain = game\_of\_life;

objects {

x\_pos : {x1,x2};  
y\_pos : {y1,y2};

};

non-fluents {

PROB\_REGENERATE = 0.9;

NEIGHBOR(x1,y1,x1,y2);  
NEIGHBOR(x1,y1,x2,y1);  
NEIGHBOR(x1,y1,x2,y2);

NEIGHBOR(x1,y2,x1,y1);  
NEIGHBOR(x1,y2,x2,y1);  
NEIGHBOR(x1,y2,x2,y2);

NEIGHBOR(x2,y1,x1,y1);  
NEIGHBOR(x2,y1,x1,y2);  
NEIGHBOR(x2,y1,x2,y2);

NEIGHBOR(x2,y2,x1,y1);  
NEIGHBOR(x2,y2,x1,y2);  
NEIGHBOR(x2,y2,x2,y1);

};

non-fluents game3x3 {

domain = game\_of\_life;

objects {

x\_pos : {x1,x2,x3};  
y\_pos : {y1,y2,y3};

};

non-fluents {

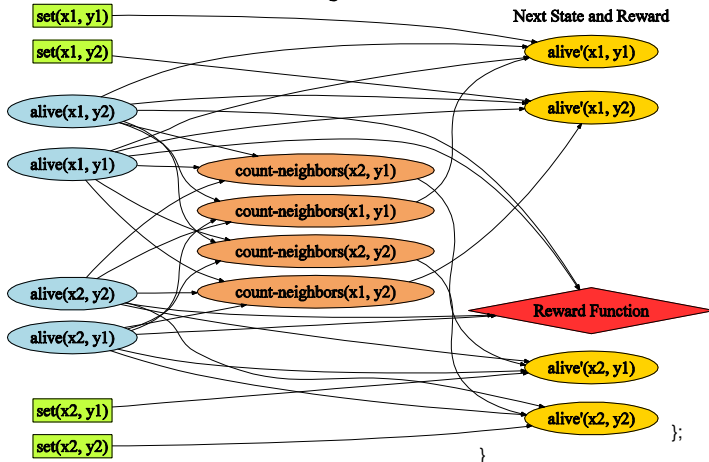
NEIGHBOR(x1,y1,x1,y2);  
NEIGHBOR(x1,y1,x2,y1);  
NEIGHBOR(x1,y1,x2,y2);  
NEIGHBOR(x1,y2,x1,y1);  
NEIGHBOR(x1,y2,x2,y1);  
NEIGHBOR(x1,y2,x2,y2);  
NEIGHBOR(x1,y2,x2,y3);  
NEIGHBOR(x1,y2,x1,y3);  
NEIGHBOR(x1,y3,x1,y2);  
NEIGHBOR(x1,y3,x2,y2);  
NEIGHBOR(x1,y3,x2,y3);  
NEIGHBOR(x2,y1,x1,y1);  
NEIGHBOR(x2,y1,x1,y2);  
NEIGHBOR(x2,y1,x2,y2);  
NEIGHBOR(x2,y1,x3,y2);  
NEIGHBOR(x2,y1,x3,y1);  
NEIGHBOR(x2,y2,x1,y1);  
NEIGHBOR(x2,y2,x1,y2);  
NEIGHBOR(x2,y2,x1,y3);  
NEIGHBOR(x2,y2,x2,y1);  
NEIGHBOR(x2,y2,x2,y3);  
NEIGHBOR(x2,y2,x3,y1);  
NEIGHBOR(x2,y2,x3,y2);  
NEIGHBOR(x2,y2,x3,y3);  
NEIGHBOR(x2,y3,x1,y3);  
NEIGHBOR(x2,y3,x1,y2);  
NEIGHBOR(x2,y3,x2,y2);  
NEIGHBOR(x2,y3,x3,y2);  
NEIGHBOR(x2,y3,x3,y3);  
NEIGHBOR(x3,y1,x2,y2);  
NEIGHBOR(x3,y1,x3,y2);  
NEIGHBOR(x3,y2,x2,y1);  
NEIGHBOR(x3,y2,x2,y2);  
NEIGHBOR(x3,y2,x2,y3);  
NEIGHBOR(x3,y2,x3,y3);  
NEIGHBOR(x3,y3,x2,y3);  
NEIGHBOR(x3,y3,x2,y2);  
NEIGHBOR(x3,y3,x3,y2);  
NEIGHBOR(x3,y3,x3,y2);

};

Current State and Actions

Intermediate @ Level 1

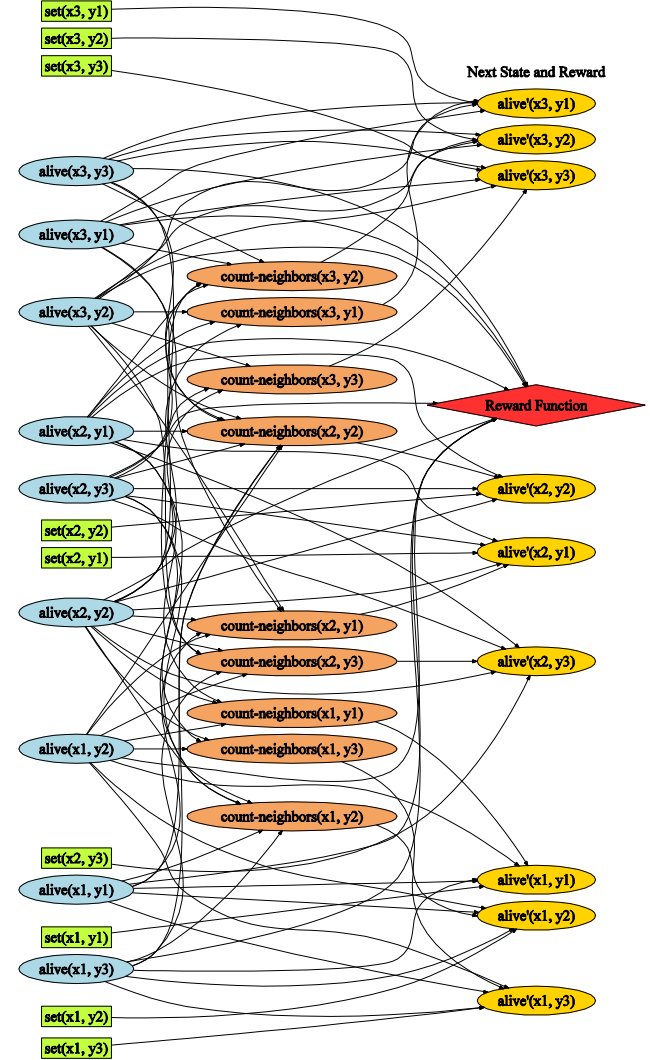
Next State and Reward



Current State and Actions

Intermediate @ Level 1

Next State and Reward




We're getting ahead of  
ourselves

Let's see how RDDDL can specify  
a binary discrete DBN+ID

# How to Represent Factored MDP?

Current State and Actions

Next State and Reward



$p$	$r$	$p'$	$P(p' p,r)$
<i>true</i>	<i>true</i>	<i>true</i>	0.9
<i>true</i>	<i>true</i>	<i>false</i>	0.1
<i>true</i>	<i>false</i>	<i>true</i>	0.3
<i>true</i>	<i>false</i>	<i>false</i>	0.7
<i>false</i>	<i>true</i>	<i>true</i>	0.3
<i>false</i>	<i>true</i>	<i>false</i>	0.7
<i>false</i>	<i>false</i>	<i>true</i>	0.3
<i>false</i>	<i>false</i>	<i>false</i>	0.7



# RDDL Equivalent

```
// Define the state and action variables (not parameterized here)
pvariables {
    p : { state-fluent, bool, default = false };
    q : { state-fluent, bool, default = false };
    r : { state-fluent, bool, default = false };
    a : { action-fluent, bool, default = false };
};

// Define the conditional probability function for each
// state variable in terms of previous state and action
cpfs {
    p' = if (p ^ r) then Bernoulli(.9) else Bernoulli(.3);

    q' = if (q ^ r) then Bernoulli(.9)
        else if (a) then Bernoulli(.3) else Bernoulli(.8);

    r' = if (~q) then KronDelta(r) else KronDelta(r <=> q);
};

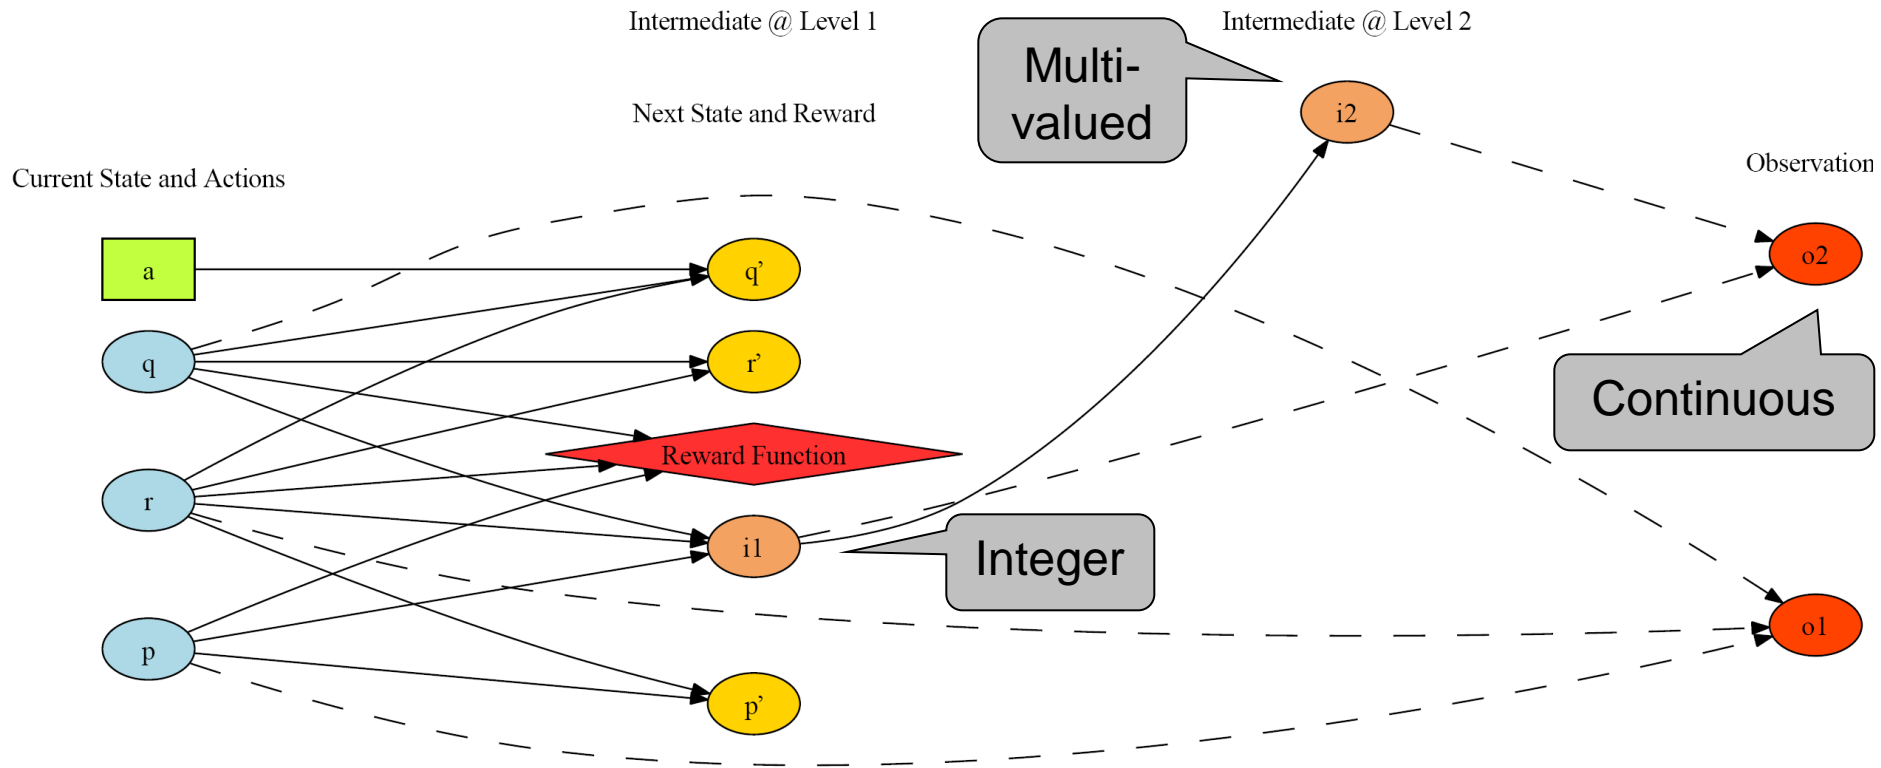
// Define the reward function; note that boolean functions are
// treated as 0/1 integers in arithmetic expressions
reward = p + q - r;
```

Can think of  
transition  
distributions  
as “*sampling  
instructions*”

# Let's look at a few more RDDDL ingredients

- enum, integer, continuous fluents
- intermediate fluents
- observation fluents (POMDP)
- more control / stochastic constructs

# A Discrete-Continuous POMDP?



# A Discrete-Continuous POMDP, Part I

```
// User-defined types
types {
    enum_level : {@low, @medium, @high}; // An enumerated type
};

pvariables {
    p : { state-fluent, bool, default = false };
    q : { state-fluent, bool, default = false };
    r : { state-fluent, bool, default = false };

    i1 : { interm-fluent, int, };
    i2 : { interm-fluent, enum_level };

    o1 : { observ-fluent, bool };
    o2 : { observ-fluent, real };

    a : { action-fluent, bool, default = false };
};

cpfs {

    // Some standard Bernoulli conditional probability tables
    p' = if (p ^ r) then Bernoulli(.9) else Bernoulli(.3);

    q' = if (q ^ r) then Bernoulli(.9)
        else if (a) then Bernoulli(.3) else Bernoulli(.8);

    // KronDelta is a delta function for a discrete argument
    r' = if (~q) then KronDelta(r) else KronDelta(r <=> q);
}
```

# A Discrete-Continuous POMDP, Part II

Integer

```
// Just set i1 to a count of true state variables
```

```
i1 = KronDelta(p + q + r);
```

Multi-valued

```
// Choose a level with given probabilities that sum to 1
```

```
i2 = Discrete(enum_level,
```

```
    @low : if (i1 >= 2) then 0.5 else 0.2,
```

```
    @medium : if (i1 >= 2) then 0.2 else 0.5,
```

```
    @high : 0.3
```

```
);
```

Real

```
// Note: Bernoulli parameter must be in [0,1]
```

```
o1 = Bernoulli( (p + q + r)/3.0 );
```

```
// Conditional linear stochastic equation
```

```
o2 = switch (i2) {
```

```
    case @low      : i1 + 1.0 + Normal(0.0, i1*i1),
```

```
    case @medium   : i1 + 2.0 + Normal(0.0, i1*i1/2.0),
```

```
    case @high     : i1 + 3.0 + Normal(0.0, i1*i1/4.0) };
```

```
};
```


Mixture of Normals

Variance comes from other previously sampled variables

# Finally: Mars Rover example

- lifting
- non-fluents
- aggregation expressions
- joint action preconditions

# Lifted Continuous MDP in RDDDL: **Simple Mars Rover**



# Simple Mars Rover: Part I

```
types { picture-point : object; };
```

```
pvariables {
```

Constant  
picture  
points,  
bounding box

```
    PICT_XPOS(picture-point) : { non-fluent, real, default = 0.0 };  
    PICT_YPOS(picture-point) : { non-fluent, real, default = 0.0 };  
    PICT_VALUE(picture-point) : { non-fluent, real, default = 1.0 };  
    PICT_ERROR_ALLOW(picture-point) : { non-fluent, real, default = 0.5 };
```

Rover position  
(only one  
rover)  
and time

```
    xPos : { state-fluent, real, default = 0.0 };  
    yPos : { state-fluent, real, default = 0.0 };  
    time : { state-fluent, real, default = 0.0 };
```

Rover  
actions

```
    xMove      : { action-fluent, real, default = 0.0 };  
    yMove      : { action-fluent, real, default = 0.0 };  
    snapPicture : { action-fluent, bool, default = false };
```

Question, how  
to make multi-  
rover?



# Simple Mars Rover: Part II

```
cpfs {
```

```
// Noisy movement update
```

```
xPos' = xPos + xMove + Normal(0.0, MOVE_VARIANCE_MULT*xMove);
```

```
yPos' = yPos + yMove + Normal(0.0, MOVE_VARIANCE_MULT*yMove);
```

```
// Time update
```

```
time' = if (snapPicture)
```

```
  then (time + 0.25)
```

```
  else (time +
```

```
    [if (xMove > 0) then xMove else -xMove] +
```

```
    [if (yMove > 0) then yMove else -yMove]);
```

Fixed time for picture

White noise, variance  
proportional to distance moved

Time proportional to  
distance moved

```
};
```

# Simple Mars Rover: Part III

// We get a reward for any picture taken within picture box error bounds  
// and the time limit.

```
reward = if (snapPicture ^ (time <= MAX_TIME))  
    then sum_{?p : picture-point} [  
        if ((abs[ PICT_XPOS(?p) – xPos] <= PICT_ERROR_ALLOW(?p))  
            ^ (abs[ PICT_YPOS(?p) – yPos] <= PICT_ERROR_ALLOW(?p)))  
        then PICT_VALUE(?p)  
        else 0.0 ]  
    else 0.0;
```

Reward for all pictures taken  
within bounding box!

**action-preconditions** {

// Cannot snap a picture and move at the same time


**snapPicture => ((xMove == 0.0) ^ (yMove == 0.0));**

};

Cannot move and take  
picture at same time.

# RDDL Recap

- Relational Dynamic Influence Diagram Language
  - Relational  
[DBN + Influence Diagram]
- Specify the probabilistic process over relations to generate next state
  - Generate “ground” DBN+ID given domain object instantiation



# RDDL Recap I

- Everything is a fluent (parameterized variable)
  - State fluents
  - Observation fluents
    - for partially observed domains
  - Action fluents
    - supports factored concurrency
  - Intermediate fluents
    - derived predicates, correlated effects, ...
  - Constant nonfluents (general constants, topology relations, ...)
- Flexible fluent types
  - Binary (predicate) fluents
  - Multi-valued (enumerated) fluents
  - Integer and continuous fluents (from PDDL 2.1)

# RDDL Recap II

- **Semantics is ground DBN + Influence Diagram**
  - Naturally supports independent exogenous events
- **General expressions in transition / reward**
  - Logical expressions ( $\wedge, \vee, \Rightarrow, \Leftrightarrow, \forall, \exists$ )
  - Arithmetic expressions ( $+, -, *, /, \Sigma_x, \Pi_x$ )
  - In/dis/equality comparison expressions ( $=, \neq, <, >, \leq, \geq$ )
  - Conditional expressions (if-then-else, switch)
  - Standard Functions:  $\text{pow}[\cdot], \log[\cdot], \text{abs}[\cdot], \text{max}[\cdot], \sin[\cdot]$
  - Basic probability distributions
    - Bernoulli, Discrete, Normal, Poisson

Logical expr.  $\{0,1\}$   
so can use in  
arithmetic expr.

# RDDL Recap III

- Goal + General (PO)MDP objectives
  - Arbitrary reward
    - goals, numerical preferences (c.f., PDDL 3.0)
  - Finite horizon
  - Discounted or undiscounted
- State/action constraints
  - Encode legal action-preconditions
    - (concurrent) action preconditions
  - Assert state-invariants
    - serve as integrity constraint checks on state
    - e.g., an elevator cannot be in two locations

# What RDDDL does not do...

- RDDDL just provides a language for specifying complex (PO)MDPs
  - For an MDP:  $\langle S, A, T, R \rangle$
  - For a POMDP:  $\langle S, A, T, R, O, Z \rangle$
- RDDDL does not define a policy
- RDDDL does not specify a planning methodology
  - It's up to **external planners** to perform planning, learning, or inference on the RDDDL domain model

# RDDL Tutorial Outline

- **Part 1: Language Overview**
  - What is probabilistic planning in PPDDL?
  - Why do we need RDDL?
  - RDDL by example
  - **Overview of RDDL solution methodologies**
- **Part 2: PyRDDL Gym**



Common question from RL crowd:  
Why RDDDL vs. a Simulator in C++?

**Answer:** Want a language that can be **compiled** into other formalisms for planning and domain analysis such as abstraction (e.g., Tensorflow / PyTorch / Jax, XADDs, Gurobi – constrained optimization).

RDDL is a disciplined subset of modern languages designed to facilitate compilation.

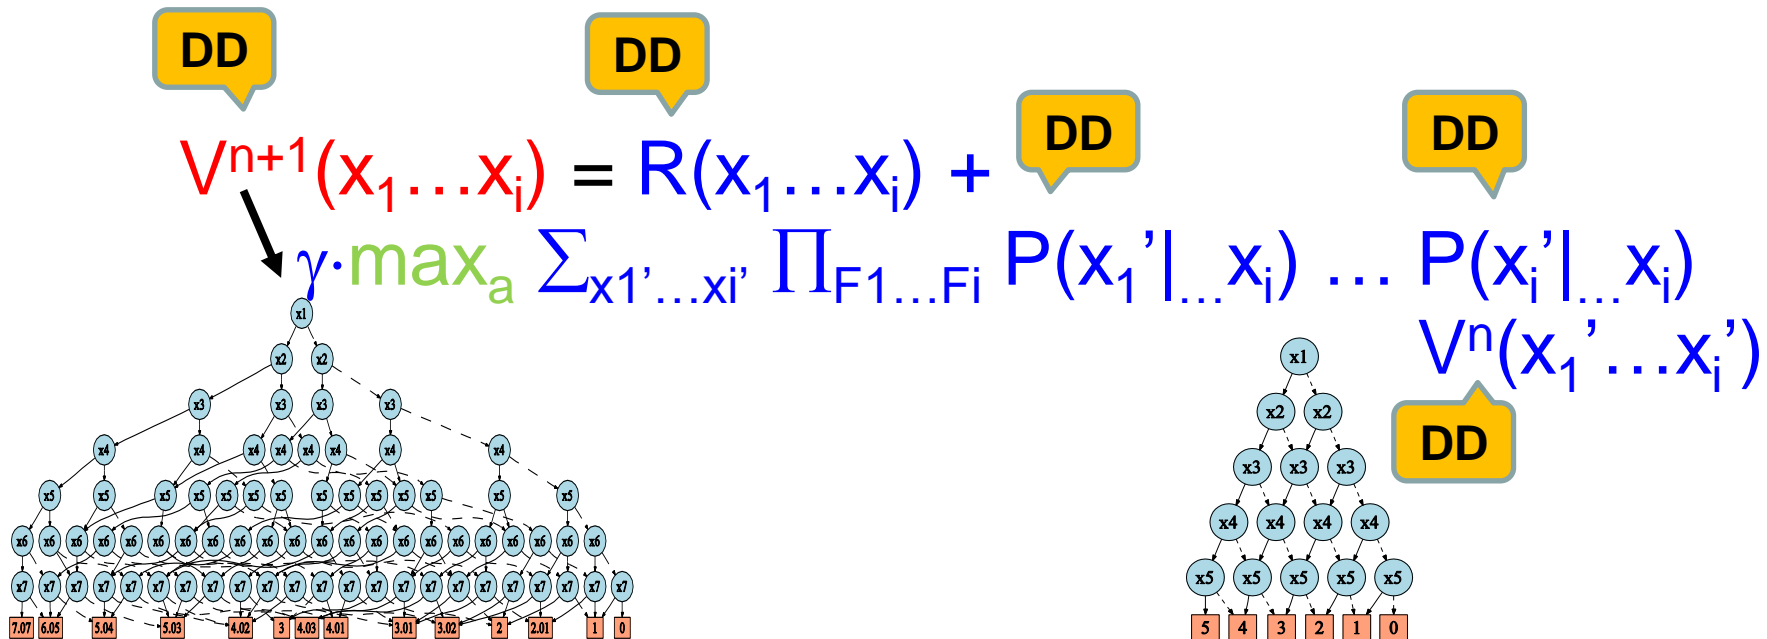
# RDDL Planning Overview

- **SOTA: compile instance to planning formalism**
  - MCTS (Discrete Search) (PROST, Keller et al, ICAPS-12) – discrete only
  - Symbolic Methods (Decision Diagrams – XADDs)
  - Planning by Backprop (Tensorplan, JaxPlan, SOGBOFA)
  - Planning by Optimization in Gurobi (Raghavan et al, AAIL-17)
- **Generalized Planning: “solve” at lifted domain level**
  - Relational / First-order MDPs (Khardon et al, Sanner et al)
  - Graph neural network policies (Symnet 1/2/3: Mausam et al)
  - Plan / policy should work for all instances

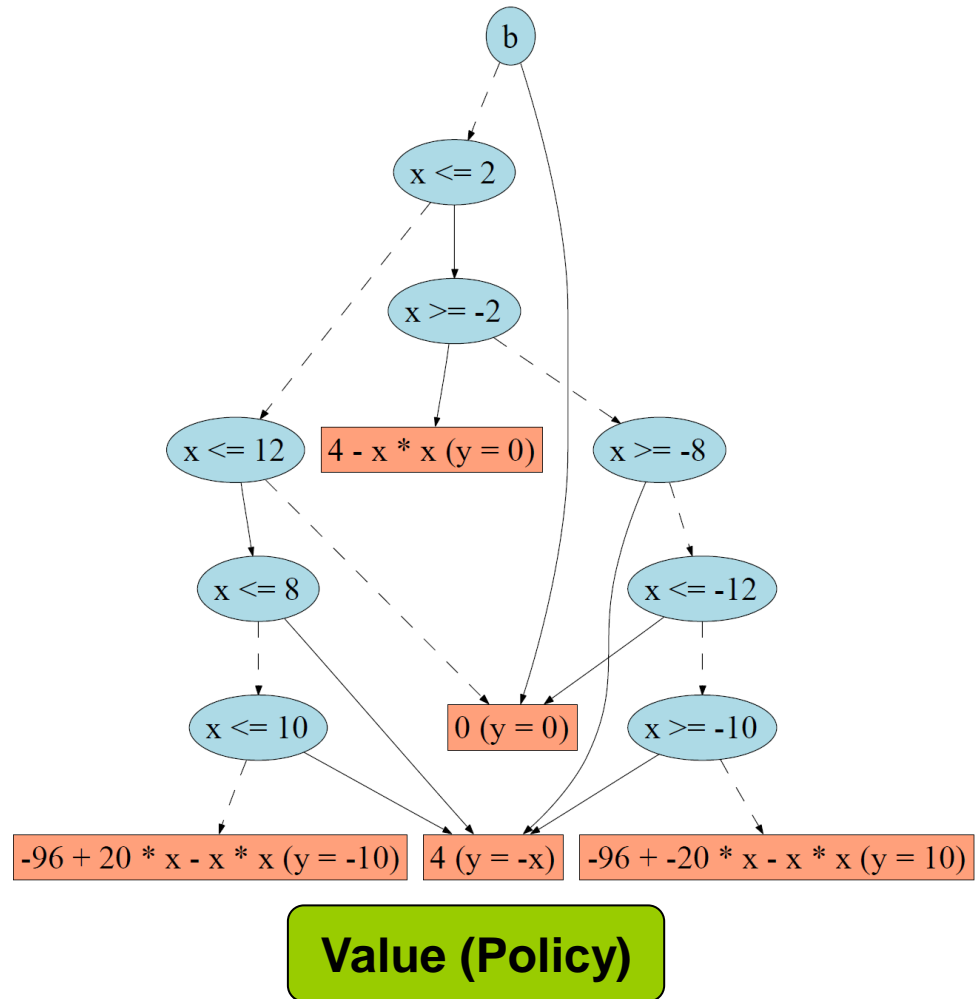




# Symbolic Decision Diagram Methods

# SPUDD for Factored MDPs

- Value Iteration using ADDs (SPUDD)
  - Can use ADDs or any DD that supports  $+$ ,  $*$ ,  $\max$
  - Bounded approximations (APRICODD)



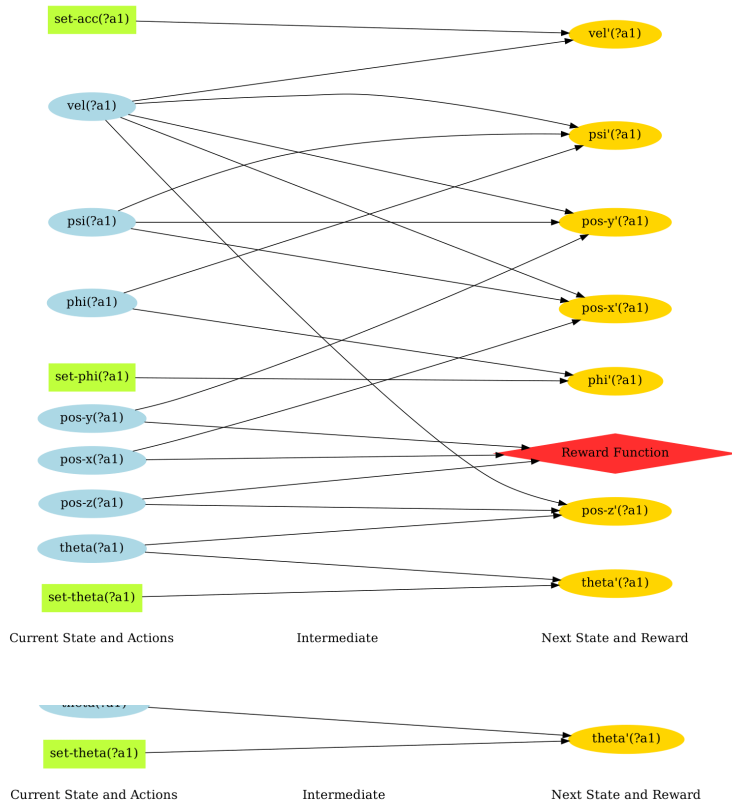
# XADDs for Discrete+Continuous MDPs



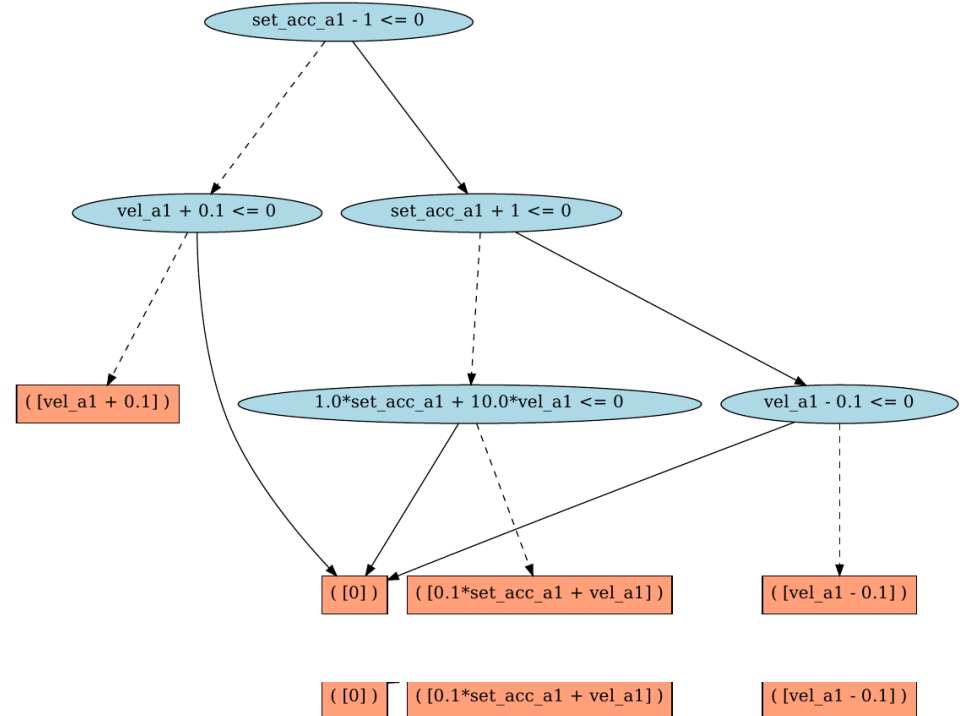
# RDDL Compiles to (X)ADDs!

- UAV Problem

DBN

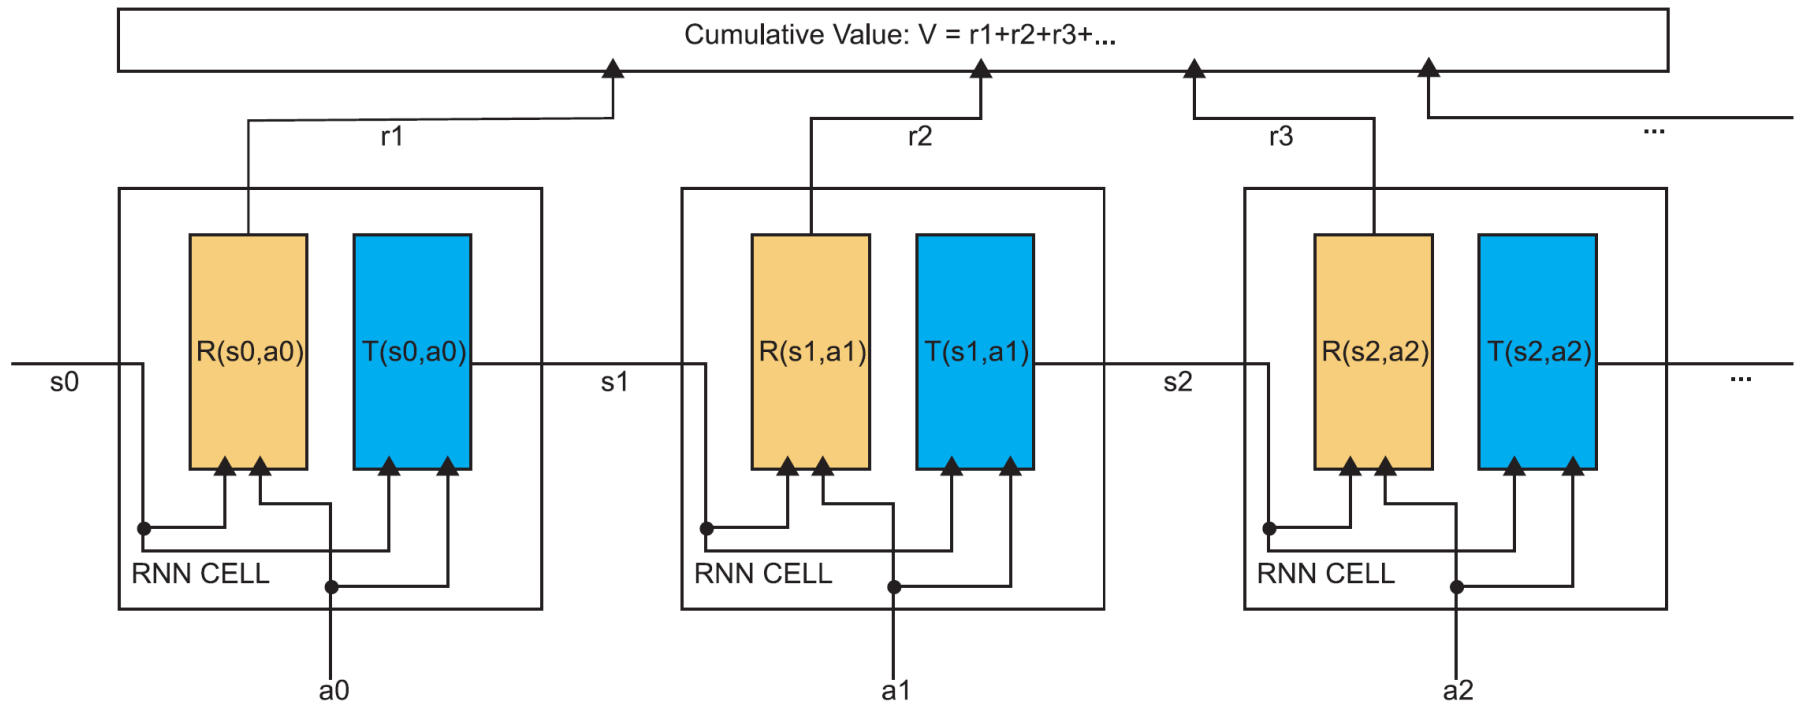


XADD for  $vel'(a1)$



# Planning by Backprop


# Tensorplan: Embed Reward and Transition in an RNN and Optimize *End-to-end*!







# GPU-based Path Planning via Tensorflow

- RMSProp makes for a great non-convex optimizer!



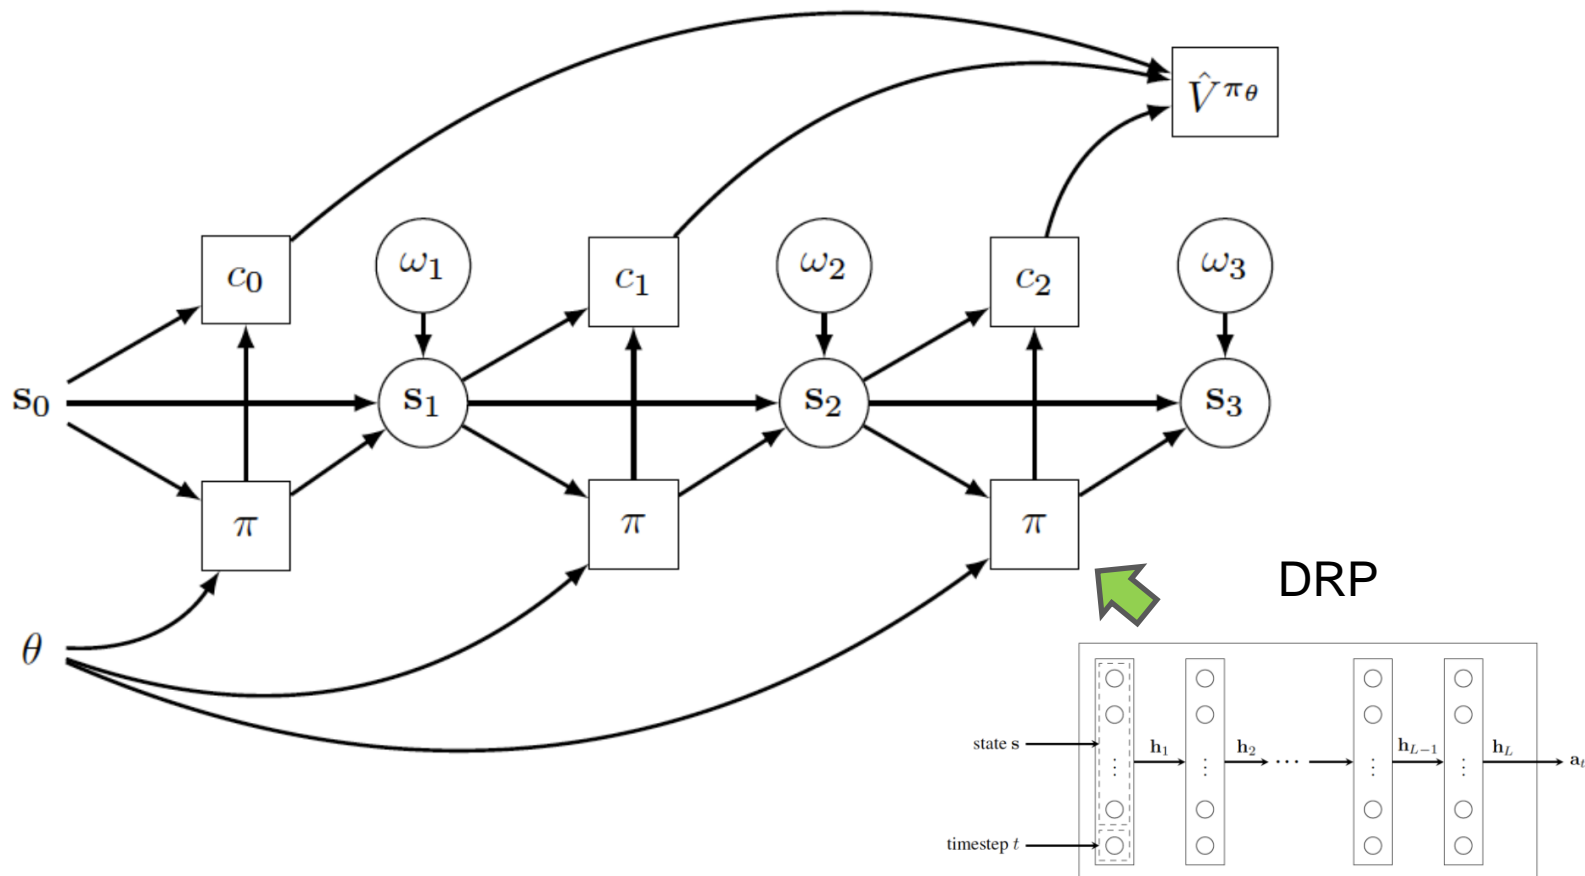
# Need Modern Non-convex Gradient Methods



RMSProp is the best-performing optimizer for planning, likely b/c it can handle piecewise structure.

# Learning Deep Reactive Policies (DRPs)

## Stochastic RNNs

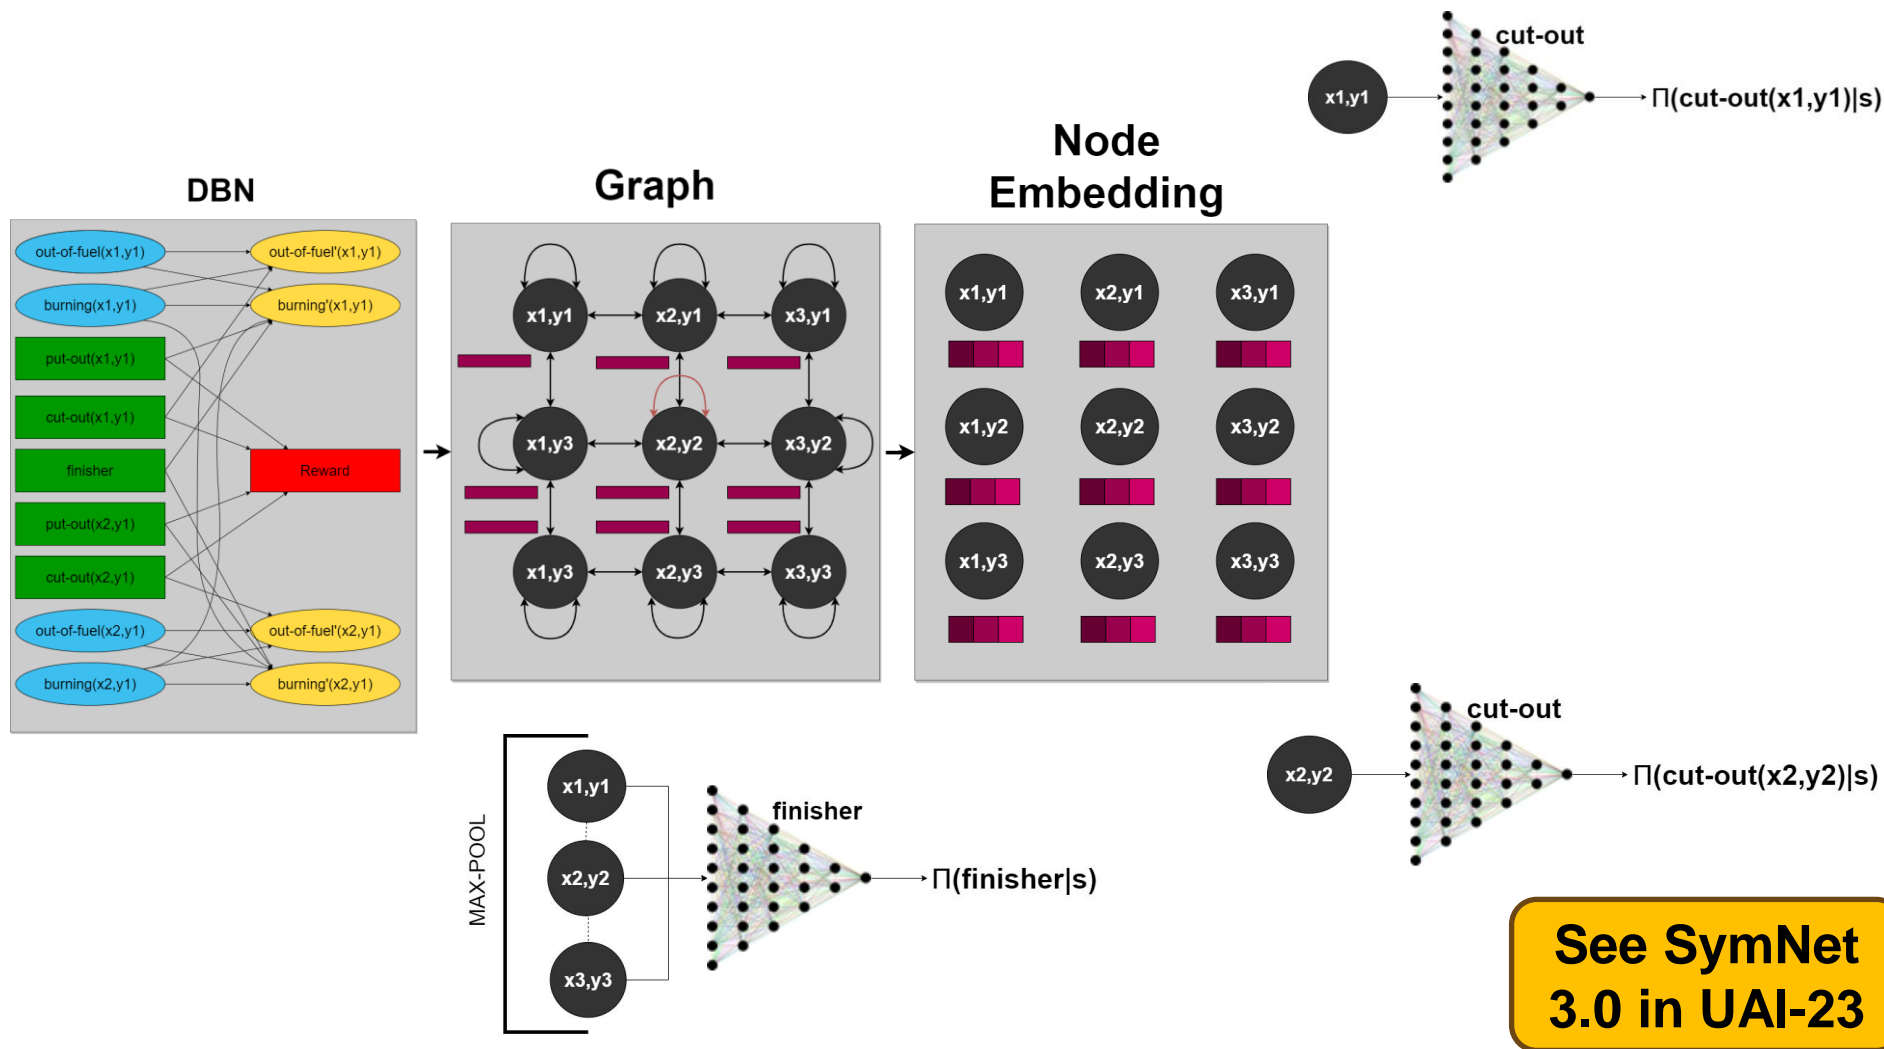


# Lifted Approaches: Generalized Planning for RDDL

SymNet (Mausam and students)

# SymNet 2.0 (Mausam et al, ICML-22)

Compile RDDDL DBN into GNN, Embed, Decode to Actions  
(GNN learning is domain instance independent)




# RDDL Tutorial Outline

- **Part 1: Language Overview**
  - What is probabilistic planning in PPDDL?
  - Why do we need RDDL?
  - RDDL by example
  - Overview of RDDL solution methodologies
- **Part 2: PyRDDL Gym**


# pyRDDLGym

Includes OpenAI Gym interface, JaxPlanner, XADDs, etc.


<https://github.com/ataitler/pyRDDLGym>



Wildfires




Mars Rover




Recommender System


**More on  
PyRDDLGym  
with worked  
examples in  
Part 2!**



Power Generation



UAVs



Elevators